

# ADvaNCE: Hashing-based Approximate DBSCAN

Tianrun Li<sup>†</sup>, Thomas Heinis<sup>¶</sup>, Wayne Luk<sup>¶</sup>

<sup>†</sup>*Tsinghua University, China*

<sup>¶</sup>*Imperial College London, United Kingdom*

## ABSTRACT

Analyzing massive amounts of data and extracting value from it has become key across different disciplines. As the amounts of data grow rapidly, however, current approaches for data analysis are no longer efficient. This is particularly true for clustering algorithms where distance calculations between pairs of points dominate overall time: the more data points in the data set, the bigger the share of time needed for distance calculations.

Crucial to the data analysis and clustering process, however, is that it is rarely straightforward: instead, parameters need to be determined and tuned first. Entirely accurate results are thus rarely needed and instead we can sacrifice little precision of the final result to accelerate the computation. In this paper we develop ADvaNCE, a new approach based on approximating DBSCAN. More specifically, we propose two measures to reduce distance calculation overhead and to consequently approximate DBSCAN: (1) locality sensitive hashing to approximate and speed up distance calculations and (2) representative point selection to reduce the number of distance calculations.

The experiments show that the resulting clustering algorithm is more scalable than the state-of-the-art as the data sets become bigger. Compared with the most recent approximation technique for DBSCAN, our approach is in general one order of magnitude faster (at most 30x in our experiments) as the size of the datasets increase.

## 1. INTRODUCTION

Unlocking the value in the masses of the data stored and available to us has become a primary concern. Medical data, banking data, shopping data and others are all analysed in great detail to find patterns, to classify behaviour or phe-

nomena and to finally predict behaviour and progression. The outcome from these analyses is ultimately hoped to predict behaviour of customers, to increase sales in marketing [8], optimise diagnostic tools in medicine to detect disease earlier, optimise medical treatments for better outcome [1, 16] etc.

Generating insightful knowledge from data, however, is not efficient today. While there exists a plethora of tools and algorithms for the analysis and clustering of data, most of them have a considerable complexity, meaning they will not scale well to the massive datasets we see today. They are very efficient on the small datasets we dealt with until recently but as the amounts of data grow rapidly these algorithms do not scale and fail to provide a fast analysis.

At the same time, data analysis rarely is a straightforward process. All clustering algorithms require the tuning of parameters, e.g., number of clusters in k-means, maximum distance ( $\epsilon$ ) in DBSCAN etc., and also the data preparation process frequently needs to be repeated, for example to adjust the level or aggregation or similar. Clearly, throughout the process of iteratively improving the analysis, the results need not be exact. Instead, by sacrificing little precision, we can substantially accelerate each analysis and thus the overall process.

More formally, the problem we address is density-based clustering, i.e., given a set of points  $P$  in  $d$ -dimensional space  $\mathbb{R}^d$  (with typically a very high  $d$ ), the goal is to find groups of points into clusters, i.e., to divide the points into dense areas separated by sparse areas. Existing approaches generally differ in how dense and sparse areas are defined but all approaches share that distance calculations account for the majority of the execution time. By reducing the number of distance calculations and approximating the remaining ones, we can trade accuracy for efficiency and scalability.

In this paper we propose ADvaNCE (Approximate DeNsity-Based ClustEring), a novel approximation approach for DBSCAN, arguably the most prominent density-based clustering approach. Our approach reduces the overhead of calculating distances with two measures. First, we use locality sensitive hashing (LSH) to approximate each distance calculation and thus accelerate it. Clearly, this step in itself already introduces approximation. Second, we approximate the dataset by only retaining representative points which summarize its structure. Doing so reduces the number of points and consequently also the number of distance calculations. As with locality sensitive hashing, also this measure approximates the result.

The resulting approach outperforms the classic DBSCAN

Permission to make digital or hard copies of all or part of this work for personal or classroom use is granted without fee provided that copies are not made or distributed for profit or commercial advantage and that copies bear this notice and the full citation on the first page. To copy otherwise, to republish, to post on servers or to redistribute to lists, requires prior specific permission and/or a fee.

Copyright 20XX ACM X-XXXXX-XX-X/XX/XX ...\$15.00.

by orders of magnitude by sacrificing the accuracy in the order of 0.001%. More importantly, ADvANCE outperforms the state-of-the-art approximate DBSCAN approaches by up to one order of magnitude (at most 30x). It scales particularly well with an increasing size of datasets for smaller  $\epsilon$ 's and consistently outperforms the state of the art.

The remainder of this paper is structured as follows. In Section 2 we discuss in detail DBSCAN along with its known optimizations but also its challenges. We then move to first give an overview of ADvANCE in Section 3 and discuss in more detail the approximate acceleration based on locality sensitive hashing in Section 4.2 and the approximation based on representative points in Section 5. To understand the benefits and limitations of ADvANCE we first analyze it analytically in Section 6 and then compare ADvANCE's efficiency to related approaches in Section 7. We then analyze its performance in detail in Section 8. We discuss the state-of-the-art of clustering algorithms (approximate and exact) in Section 9 and finally draw conclusions in Section 10.

## 2. DBSCAN AND ITS CHALLENGES

DBSCAN is arguably the most renowned density based clustering algorithm. We first give an overview of DBSCAN, present a first optimization and then motivate our approach with a first motivation experiment.

### 2.1 DBSCAN Recap

DBSCAN is a density based clustering algorithm that takes two parameters  $\epsilon$  and  $minPts$  to determine what areas of the datasets are dense/sparse and finally to compute the clustering result. For every input point  $p$ , if the neighbourhood within radius  $\epsilon$  contains at least  $minPts$  number of points, we call point  $p$  a core point.

**Definition.** Point  $p$  is density reachable from point  $q$  if there exists a sequence of points  $p_1, p_2, p_3 \dots p_n$  such that either:

1.  $p_1 = p$  and  $p_n = q$
2.  $p_1, p_2, p_3 \dots p_{n-1}$  are core points
3.  $p_{i+1}$  is in the neighbourhood of  $p_i$  with radius  $\epsilon$

Points that are not core points themselves but are density-reachable from any core point are called border points. A cluster defined by DBSCAN starts from a single core point and gradually adds all density-reachable points from any points in this cluster, which leads to a chain effect, and the cluster grows until it is finally complete. Points that are neither core points nor border points are called noise points and do not belong to any cluster.

### 2.2 Grid-based Optimization

Naive implementations of the DBSCAN algorithm [11] compute distances between all pairs of data points in the dataset in  $O(n^2)$ . A first proposed optimization uses a KD-Tree [5] to reduce the number of distance calculations between data points.

At the center of ADvANCE is a grid-based optimization [13] of DBSCAN that further reduces the distance calculations. The grid of the approach has the same dimensions as all points in the dataset. With this grid we can considerably reduce the number of distance calculations: given a point  $p$ ,

we only have to search in a fixed number of cells around  $p$  to find other points within distance  $\epsilon$  of  $p$ . More precisely, given the center cell  $c$  in which  $p$  is located, we only have to search the cells  $c'$  with  $dist(c, c') < \epsilon$ , i.e., the neighbouring cells  $N(c)$  (red in Figure 1).

The grid uses uniform spacing of the cells in all dimensions, making assignment of points to cells and related calculations straightforward. With the grid, the algorithm has four distinct phases.

#### Grid Construction:

In the assignment phase the algorithm iterates over all points and maps each point to the grid cell that encloses it. To eventually reduce the number of distance calculations, we choose the grid cell width as  $\frac{\epsilon}{\sqrt{D}}$

with  $D$  as the dimension, guaranteeing that all points within a cell are by definition within distance  $\epsilon$  of each other and thereby form a cluster.

**Determining Core Points:** The goal of this phase is to determine core points. For this phase the choice of grid cell width used in the first phase is crucial. As mentioned previously, if the grid cell width in the first phase is set to  $\frac{\epsilon}{\sqrt{D}}$ , then the diagonal of the cell is  $\epsilon$  and consequently the distance between any two points in the same cell is at most  $\epsilon$ . If a cell  $c$  contains more than  $minPts$  points, all the points inside cell  $c$  are consequently core points.

**Merge Clusters:** Most important in this phase is that we consider each non-empty cell as a small cluster. Further, if two core points in two different cell are within distance less than  $\epsilon$  of each other, these two points belong to the same cluster and the clusters need to be merged.

To compute all clusters, the optimization generates a graph  $G = (V, E)$  where  $V$  represents all non-empty cells. Given two cells  $c_1$  and  $c_2$ ,  $E$  contains an edge between  $c_1$  and  $c_2$  if there exists core points  $p_1$  in  $c_1$  and  $p_2$  in  $c_2$  with  $dist(p_1, p_2) < \epsilon$ . The result of the merge cluster phase consequently are the connected components of graph  $G$  representing the clusters.

**Determine Border Points:** Given a non-core point  $p$  in cell  $c$  (by definition a cell with less than  $minPts$ ), we check all core points in neighbouring cells  $N(c)$ , and find the core point  $q$  with the minimum distance to  $p$ . Then  $p$  is a border point and belongs to the cluster of core point  $q$ . If there are no core points in  $N(c)$ , then  $p$  is a noise point and does not belong to any cluster.

### 2.3 DBSCAN Challenges

A considerable challenge of DBSCAN (and other clustering approaches) is the number of distance calculations and the time needed for them. Even for the optimized DBSCAN algorithm [13] the number of distance calculations is substantial. To make matters worse, the number grows rapidly with increasing dataset size.

We demonstrate this problem with a motivation experi-





Figure 1: Neighboring cells of  $c$  are in red.

ment where we increase the size of the dataset from 0.125 millions to 1 million. The data sets used have five dimensions and are randomly generated [12]. For DBSCAN we set the parameters to  $\epsilon$  to 5000 and the minimum number of points to 100 [12]. As the results in Figure 2 show, the share of total time needed to run DBSCAN starts with a few seconds for a very small dataset of 125'000 it grows rapidly to a substantial share of more than 60% for a dataset of 1 million. We have repeated the experiment for each dataset size three times and report the average.

Clearly, the share of the overall time spent on distance calculations is substantial and grows rapidly. As we will show in the remainder of this paper, we speedup DBSCAN considerably and enable it to scale by using approximation in two respects. First we use approximation to reduce the total number of distance calculations and second we use approximation to accelerate the remaining distance calculations. As we will show, only little precision has to be sacrificed using ADvANCE to accelerate the clustering process considerably.



**Figure 2: Share of overall time needed for distance calculations.**

### 3. ADVANCE OVERVIEW

ADvANCE uses the key insight that for large multi-dimensional datasets, the time needed to find clusters is primarily driven by the number of distance computations needed. While previous work [13] has substantially reduced this number, it is still considerable and grows rapidly for increasingly big datasets. By approximating the cluster computation in a grid-based DBSCAN [13] we considerably accelerate the process.

The first approximation reduces the the number of distance calculations, or more precisely, the number of neighboring cells a cell has to speed up clustering. Using locality sensitive hashing [2] (LSH) we can approximate distances between the points and efficiently establish whether two points are within distance  $\epsilon$  of each other.

A second approximation reduces the points considered in each cell. To decide whether or not two grid cells, each containing more than  $minPts$  and thus containing only core points, should be joined, not all points need to be considered. To obtain an approximate result it is sufficient to take into account primarily points closer to the border of the cell: if each of two cells have a point within distance  $\epsilon$  of each other then there is no point to test further points.

Particularly, the former optimization merges the cells with at least  $minPts$  (which are consequentl all core points) approximately and efficiently. It does, however, only join a

subset of neighboring cells and consequently several iterations are needed to join all neighboring cells of core points. ADvANCE consequently is run iteratively to join cells. It runs iteratively until the results converges, i.e., until the result no longer changes between several iterations (set as a parameter).

## 4. ADVANCE: APPROXIMATE NEIGHBORHOOD

While the basic Grid based algorithm is very efficient in two dimensions with  $D = 2$  and in case the width of each cell is set to  $\frac{\sqrt{D}\epsilon}{2}$  (and the diagonal of each cell consequently is  $\epsilon$ ), in higher dimensions the number of  $N(c)$  cells grows considerably. For example, in five dimensions the cell width becomes very small with  $\epsilon/\sqrt{5}$ , meaning that three cells need to be searched in each dimension and  $N(c)$  contains  $7^5$  cells. Merging cells into clusters will be very time consuming as all of  $N(c)$  ( $7^5$  cells in five dimensions) has to be searched for every core point.

In the following we first discuss the basic concept of approximating the neighborhood and then present how to design the algorithm that (a) approximates the neighborhood and (b) clusters based on the approximation.

### 4.1 Using Locality Sensitive Hashing to Approximate

As is discussed previously, searching in the neighbouring cells  $N(c)$  is crucial to the execution of grid-based DBSCAN [13] as it is an integral part of every step (except when constructing the grid). The number of cells in  $N(c)$ , however, is growing uncontrollably with an increasing number of dimensions. The basic idea of our approximate algorithm is to find an approximate neighbour for each core point using locality sensitive hashing instead of searching in the potentially large number of cells in  $N(c)$ . We consequently use the approximate neighbours to finish DBSCAN with a controllable precision.

In the context of ADvANCE we consequently use locality sensitive hashing (LSH [2]) to approximate and reduce the neighborhood that needs to be considered. The LSH functions  $H$  we use are random hyperplane projection functions. More precisely, given two points  $p, q$  in dimension  $D_1$  (the dimension of the input data) we use  $H$  a set of random hyperplane distance functions to project  $p$  and  $q$  into dimension  $D_2$ , thereby approximating the distance between them. With this, we can assert the following about the approximate distance between  $p$  and  $q$ :

1. if  $dist(p, q) < \epsilon$ , then for every function  $H_i \in H$ , we have  $H_i(p) - H_i(q) < \epsilon$
2. if  $dist(p, q) > \epsilon$ , then for every function  $H_i \in H$ , there is a considerable chance that  $H_i(p) - H_i(q) > \epsilon$  and there is a small possibility that  $H_i(p) - H_i(q) < \epsilon$

With this we consequently can construct a new grid  $NG$  in  $D_2$  with  $cellWidth = \epsilon$  and assign each point to the corresponding cell. Given a point  $p$ , we define the cell  $c$  in the new grid  $NG$  that contains  $p$  as the approximate neighbour of  $p$ , denoted by  $AN(p)$  and all points in  $c$  will serve as approximate neighbour points of  $p$ . The algorithm is shown in pseudocode in Algorithm 1.

Given our approach to approximation, it is possible, that points that are within the  $\epsilon$  neighbour sphere  $B(p, \epsilon)$  can not

---

**Algorithm 1:** HashAndAssign

---

**Input:**  $P$ : set of input points  
 $\epsilon$ : distance threshold  
**Output:**  $NG$ : grid in higher dimension  
**Data:**  $D_1$ : dimension of input data  
 $D_2$ : dimension to project to

Create  $D_2$  random hyperplane in  $D_1$  space  
 Initialise uniform grid  $NG$  in  $D_2$  space with cellWidth  
 $= \epsilon$

```

foreach  $Point\ p \in P$  do
  foreach  $i \in D_2$  do
     $dist[i] = \text{distance between } p \text{ and hyperplane}[i]$ 
    set coordinates of  $p$  in  $D_2$  to  $dist$ 
  assign  $p$  with coordinates  $dist$  to  $NG$ 
return ( $NG$ )

```

---

be found in a neighbour cell in  $NG$  due to the approximation. In the grid  $NG$  in  $D_2$  space  $B(p, \epsilon)$  may be split into two parts. We address this issue by iterating and accumulating the hashing results to obtain a good approximation of the true  $\epsilon$  neighbours.

## 4.2 ADvANCE-LSH - Approximate, LSH-based DBSCAN

Using the approximate neighbours in  $NG$  generated based on locality sensitive hashing we propose ADvANCE-LSH, our approximate grid-based DBSCAN algorithm using LSH. The functions we use are hashing functions with a p-stable distribution [10]. Independent of the dimensionality of the input data, we project to 8 dimensional space. An overview in pseudocode is given in Algorithm 2.

---

**Algorithm 2:** ADvANCE-LSH - approximate, LSH-based DBSCAN

---

**Input:**  $P$ : set of input points  
 $\epsilon$ : distance threshold  
 $minPts$ : minimum number of points to form a cluster

$G = \text{ConstructGrid}(P)$   
 $\text{DetermineCorePoints}(G)$

```

while true do
   $NG = \text{HashAndAssign}(P);$ 
   $\text{DetermineCorePointsHash}(NG, G, P);$ 
   $\text{MergeClustersHash}(NG, G, P);$ 
   $\text{DetermineBorderPointsHash}(NG, G, P);$ 

```

---

The *Construct Grid* and *Determine Core Points* functions are very similar to the functions explained for the basic grid-based DBSCAN implementation. In this version we support higher dimensions. Furthermore in *Determine Core Points* we only mark the points as core points when the cell contains more than  $minPts$  points and leave it to *Determine Core Points Hash* to determine which of the remaining points are core points.

The condition for the while loop in Algorithm 2 can be set

to a fixed number of iterations. Alternatively, and to control the precision of the result better, we can use a heuristic termination condition that counts the number of core points or the number of clusters and stops when the accuracy no longer improves.

The function *Determine Core Points Hash* and *Determine Border Points Hash* are the approximate versions of *Determine Core Points* and *DetermineBorderPoints* described previously. More precisely, in *Determine Core Points Hash*, for every non-core point  $p$ , we count the number of points in  $AN(p)$  within  $\epsilon$  of point  $p$ . Although  $AN(p)$  is not equal to  $B(p, \epsilon)$ , we can converge to it through iterations. The accumulation of  $AN(p)$  in successive iterations will converge to  $B(p, \epsilon)$  and we can gradually compute the full result. In *Determine Border Points Hash* we find the nearest core point for every non-core points  $p$  in  $AN(p)$  generated by LSH. This step can also be repeated several times to improve the accuracy.

The function *Merge Clusters Hash* is the approximate version of *Merge Clusters* discussed previously. If two core points  $p$  and  $q$  are in the same cell in  $NG$ , and their distance is less than  $\epsilon$  the two small clusters that contain  $p$  and  $q$  are merged. The algorithm is illustrated in Algorithm 3.

---

**Algorithm 3:** Approximate Cluster Merging

---

**Input:**  $NG$ : grid in higher dimension  
 $G$ : initial grid  
 $P$ : set of points

```

foreach  $Point\ p \in P$  do
   $c \leftarrow \text{cell in } NG \text{ containing } p$ 
  foreach  $Point\ q \in c$  do
    if  $dist(q, p) < \epsilon$  then
       $c_1 \leftarrow \text{cell in } G \text{ containing } p$ 
       $c_2 \leftarrow \text{cell in } G \text{ containing } q$ 
      merge  $c_1$  and  $c_2$ 
      break

```

---

In Algorithm 3 we merge the clusters in  $D_1$  using the hashing result in  $D_2$  space. Crucially, we add a break in the for-loop because we do not merge all possible clusters in one go but instead use a number of iterations to gradually merge the clusters and terminate when we meet the condition of the while loop in the main algorithm (Algorithm 2).

## 5. ADVANCE: REPRESENTATIVE POINTS APPROXIMATION

Given the basic approximation discussed before, we can further accelerate Grid-based DBSCAN featuring locality sensitive hashing called ADvANCE-LSH. More particularly, in case of very dense datasets, the number of points in a cell that need to be compared to points in neighbouring cells is considerable, leading to substantial overhead. This leads to a bottleneck in the main iteration described previously as we have to iterate through all points in a cell and need to determine their relationship to points in neighbouring cells.

To address this issue we have designed a further optimisation. Given  $minPts$ , which is the minimum number points to form a cluster, we set the maximum number of points in each cell to be  $maxNum$  where  $maxNum > minPts$  and

ignore the other points in the main iteration.

For every cell  $c$  in the original grid  $G$  there are three cases:

1. if  $|c| < \text{minPts}$  then we can not (yet) determine whether the points in  $c$  are core points and  $c$  is thus not affected by this approximation
2. if  $|c| \geq \text{minPts}$  and  $|c| \leq \text{maxNum}$  then all points in  $c$  are core points and this cell will also not be affected by this approximation
3. if  $|c| > \text{maxNum}$  then it follows that  $|c| > \text{minPts}$  and all points in  $c$  are core points. All points in  $c$  belong to the same cluster so we only need to identify which cluster this cell  $c$  belongs to during the main iteration. Although we ignore  $|c| - \text{maxNum}$  points we can still determine which cluster they belong to.

Key to the approach is to choose  $\text{maxNum}$  points such that we do not lose information (or at least can curb the information lost).

Using Algorithm 4 we can select most (up to  $\text{maxNum}$ ) of the points that are near the border of each cell and in the *Merge Clusters Hash* function we calculate the distance to determine the relationship between different cells. The points on the border of each cell are the most useful in determining what cells should be merged while the points near the center of the cell are not very useful and we potentially ignore them. Consequently Algorithm 4 will first calculate the geometric center of all points and take the first  $\text{maxNum}$  points furthest away from the geometric center.

---

**Algorithm 4:** Reducing Points in a Cell

---

**Input:**  $G$ : initial grid

$\text{maxNum}$ : maximum number of points per cell


```

foreach Cell  $c \in G$  do
     $gc \leftarrow$  geometric center of all points in  $c$ 
    foreach Point  $p \in c$  do
         $diff[p] \leftarrow dist(p, gc)$ ;
    sort(diff);
    take first  $\text{maxNum}$  from diff

```

---

To illustrate the algorithm we randomly generate 400 points in 4 cells and set the  $\text{maxNum}$  in each cell to be 50 and obtain the result in Figure 3: the red-small points in the center of each cell are points that are ignored by *Merge Clusters Hash* algorithm while the big blue points around the border are still considered. Even if we, as in this example, ignore half of the points in each cell the relationship between each cell can remain correct.



**Figure 3:** Neighboring cells of  $c$  in red.

## 6. ADVANCE: ANALYTICAL ANALYSIS

To use ADvANCE it is crucial to appreciate its benefits and to understand its limitations. In the following we therefore first analyze the accuracy of the algorithm and then reason about its time complexity.

### 6.1 Algorithm Accuracy

As ADvANCE only approximates the DBSCAN result, it is of course crucial to analyze its accuracy.

In the functions *DetermineCorePointsHash* and *Merge Clusters Hash* we use the approximate neighbors  $AN(p)$  based on locality sensitive hashing instead of the actual  $\epsilon$  neighbors  $B(p, \epsilon)$ . This may lead to a decreased accuracy in the following scenarios:

1. Given a point  $p$ , the number of points in  $AN(p)$  is less than  $\text{minPts}$ , but  $B(p, \epsilon)$  has more than  $\text{minPts}$  points. A core point  $p$  may therefore not be identified as core point.
2. Given two cells  $c_1$  and  $c_2$ , two core points  $p_1 \in c_1$ ,  $p_2 \in c_2$  and  $dist(p_1, p_2) < \epsilon$ , but  $p_1$  and  $p_2$  are not in the same bucket after locality sensitive hashing. The two cells that should be merged may not be merged.

In both scenarios, merging of clusters or identification of core points may not occur. Crucial to ADvANCE, however, is that it merges clusters in multiple iterations. As the number of iterations grows ADvANCE will gradually determine more core points and merge more clusters that should be merged. The impact of these two scenarios on the accuracy will consequently decrease monotonically and the algorithm will converge depending on the stop criteria which is either a fixed number of iterations or no change (in clusters or core points) in the last iterations.

Also note that, we still perform distance calculations for points in  $AN(p)$  to select the points that are truly in  $B(p, \epsilon)$  and use these points to determine core points or merge clusters. We can therefore rule out the following two cases to occur in ADvANCE:

1. Point  $p$  is determined as core point but actually it is not.
2. Two cells are merged together but should not be.

In conclusion, the accuracy of ADvANCE will increase monotonically with the number of iterations and an excessive number of iteration therefore cannot adversely affect the accuracy.

### 6.2 Time Complexity

The most time consuming part of ADvANCE is the iterations of hashing and merging. In the following we thus analyse the time complexity of each iteration of ADvANCE. Since *DetermineCorePointsHash* and *MergeClustersHash* perform a linear search in the approximate neighborhood  $AN(p)$ , the expected size of the  $AN$ , which is also the expected number of collisions for each query point in locality sensitive hashing, is the key to determining the running time of each iteration. We first discuss preliminaries, i.e., characteristics of level- $i$  cells (number and distance between them) and the  $p$ -stable hashing functions used (probability of collision) before we reason about the time complexity.

### Level-i Cell Characteristics:

Let  $P$  be the set of points in  $\mathbb{R}^d$  where  $d$  is a constant (the number of dimensions), and we construct the original grid  $G$  for all points in  $P$ . We start by defining *level-i cells*.

#### Definition. Level-i cells:

For a point  $q$  we find the cell  $c$  in  $G$  that contains  $q$ .

Cells that are directly connected to  $c$  are level-0 cells.

For level- $i$  cells, the outer cells that directly connected to level- $i$  cells are level- $(i+1)$  cells.

An example of level- $i$  cells in a 2D grid is showed in Figure 4. Relative to the red cell, the yellow cells are level-0 cells, the green cells are level-1 cells, and so on.

The number of level- $i$  cells  $N_i$  is

$$N_i = (2i + 3)^d - (2i + 1)^d = O(i^{d-1})$$

For every point  $q$  in level- $i$  cells, the minimum distance between  $q$  and  $p$   $M_i$  is

$$M_i = i \times \text{cellWidth} = i \times \frac{\epsilon}{\sqrt{d}}$$

The maximum number of points in level- $i$  cell is a constant  $\text{maxNum}$  as we discussed in the context of representative points approximation.

### Hashing Functions Characteristics:

If  $f_p(t)$  is the probability density function of the absolute value of the p-stable distribution then for two vectors  $v_1, v_2$  and with  $c = \|v_1 - v_2\|$ , the probability of collision  $p(c)$ .

$$p(c) = \Pr_{a,b}[h_{a,b}(v_1) = h_{a,b}(v_2)] = \int_0^r \frac{1}{c} f_p\left(\frac{t}{c}\right) \left(1 - \frac{t}{r}\right) dt$$


where  $a, b$  and  $r$  are parameters for p-stable LSH functions. In order to get a desirable hash result, we can also concatenate  $k$  functions  $h$  in  $H$  and redefine the hash function as  $g(v) = h_1(v), h_2(v) \dots h_k(v)$ . In this case, the probability of collision  $p(c)$  is as follows:

$$p^k(c) = \frac{1}{c^k} \left[ \int_0^r \frac{1}{c} f_p\left(\frac{t}{c}\right) \left(1 - \frac{t}{r}\right) dt \right]^k$$

With  $p$  (the possibility function defined by a p-stable LSH) and  $k$  (the number of hash functions used) the number of collisions for a query point  $q$  then is:

$$E[\#Collision] = \sum_{x \in D} p^k(\|q - x\|)$$

where  $p^k(\|q - x\|)$  is the possibility of collision if two points with distance of  $\|q - x\|$  and  $D$  is the size of the dataset.



**Figure 4: Cells neighboring the red cell are level-0 cells.**

### Time Complexity Analysis:

Since we divide the whole dataset into cells of levels, we can redefine the number of collision by level- $i$  cells. In order to determine an upper bound for the number of collisions, we can assume that all points in level- $i$  cells have the distance  $M_i$ , which is the minimum distance they can have.

$$E[\#Collision] = \sum_{x \in D} p^k(\|q - x\|)$$

$$\leq \sum_{i=0}^{\infty} \text{maxNum} \times N_i \times p^k(M_i)$$

$$\leq \sum_{i=0}^{\infty} \text{maxNum} \times O(i^{d-1}) \times p^k\left(\frac{i\epsilon}{\sqrt{d}}\right)$$

$$\leq \sum_{i=0}^{\infty} \text{maxNum} \times O(i^{d-1}) \times \left(\frac{\sqrt{d}}{i\epsilon}\right)^k$$

$$= \text{maxNum} \times \left(\frac{\sqrt{d}}{\epsilon}\right)^k \times \sum_{i=0}^{\infty} \frac{1}{O(i^{k-d+1})}$$

The expected number of collisions is proportional to an infinite series of  $i$ . Given any  $d$  (the dimension of the input data) we can always find  $k \geq (d+1)$  and the infinite series of  $i$  will consequently converge to a constant regardless of how big  $i$  is. The expected number of collision can be expressed by the following function, where  $C$  is a constant independent of the size of the dataset.

$$= \text{maxNum} \times \left(\frac{\sqrt{d}}{\epsilon}\right)^k \times C$$

The expected size of the approximate neighbor  $AN(p)$  therefore is  $O(1)$ . In function *Determine Core Point Hash* and *Merge Clusters Hash*, we perform a linear search in  $AN$  for every point, making the time complexity  $O(n)$ . Also, the hashing process itself is  $O(n)$  and so the time complexity of one whole iteration consequently is  $O(n)$ .

## 7. EXPERIMENTAL EVALUATION

In this section we first describe the experimental setup and then we perform a sensitivity analysis using synthetic datasets to understand the performance by changing one workload variable at a time. To demonstrate the benefits of our approach we also use real world datasets. Before we conclude we also analyze our algorithm in depth by discussing a breakdown of its performance.

### 7.1 Experimental Setup

The experiments are run on a Linux Ubuntu 2.6 machine equipped with Intel(R) Xeon(R) CPU E5-2640 0 CPUs running at 2.5GHz, with 64kb L1, 256KB L2 and 12MB L3 cache and 8GB RAM at 1333MHz. The storage consists of 2 SAS disks of 300GB capacity each. Storage, however, is only used once to read the data into memory. In the following we discuss the software setup as well as the configuration details.



## 7.2 Software Setup & Datasets

We compare our algorithm against the most recent approximate DBSCAN implementation,  $\rho$ -approximate DBSCAN available and set  $\rho$  to 0.001 as recommended [12]. We use two different versions of our approach, the first using the hashing approximation only (ADvaNCE-LSH) and the second using the point reduction as well (ADvaNCE). Our approach is configured to use eight p-stable hashing functions and convergence stops once no more clusters are merged after 10 iterations.

Each algorithm implemented uses a single CPU core to ensure a fair comparison. Our approach is written in C++ while we use the executable provided for [12]. All implementations are compiled using g++ with -O3 turned on.

All experiments are repeated five times and the average execution time is reported.

We use the same synthetic datasets as in [12] defined with a dimensionality  $d$ , a restart probability  $\rho_{restart}$ , a target cardinality and a noise percentage  $\rho_{noise}$ . The datasets are generated using a random walk: the walk moves step by step a distance of  $r_{shift}$  towards a random direction every  $c$  number of steps. In each step, with probability  $\rho_{restart}$ , the walker restarts by jumping to a random location in the data space, resetting the step counter to  $c$ . No matter if a restart has happened, the walker produces a point uniformly at random in the ball centered at its current location with radius 100. We repeat a total of  $n \times (1 - \rho_{noise})$  steps generating the same number of points and add  $n \times \rho_{noise}$  uniformly distributed noise points. Put simply, we start to generate a cluster and in every step, with probability  $\rho_{restart}$ , start to generate a new cluster.

We also use two real datasets [4] the *Household* dataset in 7D as well as the *KDD Cup '99* dataset (after PCA on the attributes) in 9D. The household dataset measures electric power consumption in one household with a one-minute sampling rate over a period of almost 4 years resulting in 2'049'280 data points. Points in the original database with missing coordinates were removed. The KDD Cup '99 dataset consists of 4'898'431 data points representing detected network intrusion events.

All data, synthetic and real, is normalized in a domain of  $[0, 10'000]$  in all dimensions.

## 7.3 Accuracy Metric

To measure the accuracy of the approximate result we compare it to the exact result. We use the popular omega-index to measure the quality of the clustering result [9, 14]. The omega-index assess the similarity of two clusterings as the ratio of consistent pairs of points, i.e., a point  $p$  needs to be in the same clusters in both clustering solutions. We compare the exact result obtained with the grid-based DBSCAN algorithm with either version of ADvaNCE. A value of the omega-index of 100% consequently means ADvaNCE's result is precise and all values below that indicate how much the approximate result deviates from the precise result.

## 7.4 Synthetic Data

The synthetic datasets we use are generated as described previously. The DBSCAN parameter  $minPts$  is set to 100 [12].  $maxNum$  for the representative points approximation is set to  $\sqrt{minPts \times M}$  where  $M$  is the maximum points per cell.

### 7.4.1 Increasing Dataset Size

In this first experiment we compare the execution time of our approach (both optimisations) with the most recent DBSCAN approximation technique. We cannot compare it with a basic, accurate DBSCAN as this takes far too long to execute. We execute the experiment in 5, 7, and 9 dimensions and increase the number of data points from 100'000 to 10 millions.  $\epsilon$  is set to 5000.

As can be seen in Figure 5, our approach using both approximations (ADvaNCE) outperforms  $\rho$ -approximate DBSCAN [12] consistently and almost up to two orders of magnitude (for 9 dimensions). As the number of dimensions increases, the execution time of  $\rho$ -approximate DBSCAN grows. Our approach on the other hand shows consistent performance and indeed improves execution time with an increase in the number of dimensions. This is because approximation thorough hashing works particularly well in higher dimensions. Overall our approach scales better with an increasing number of points in the dataset.

The gap between the two versions of our approach, ADvaNCE-LSH and ADvaNCE narrows as the number of dimensions grows. This effect can be attributed to hash-based approximation improving with an increasing dimensionality. As a consequence, the relative contribution of the representative points approximation shrinks and does not improve the overall results considerably for higher dimensions.

In Figure 6 we measure the number of distance calculations instead of the execution time for the Grid-based DBSCAN implementation as well as both versions of ADvaNCE (we cannot measure it for  $\rho$ -approximate DBSCAN as we only have access to the executable). Given the considerable execution time of analysing and measuring the number of distance calculations, we only measure the number of distance calculations from 0.2 to 1.6 million data points. Comparing the trends with Figure 5, we see the similar shape the number of distance calculations as well as the execution time has, clearly indicating that the number of distance calculations is indeed the major contributor to the overall execution time.

### 7.4.2 Increasing Epsilon

In a next experiment we compare the approaches with increasing  $\epsilon$  ranging from 5'000 to 50'000. Also here the basic, precise DBSCAN version takes too long to execute, we cannot include it in the results and consequently only compared ADvaNCE with  $\rho$ -approximate DBSCAN.

As the results in Figure 7 show, for an increasing  $\epsilon$ , the performance of  $\rho$ -approximate DBSCAN improves at first but then grows again (as can be seen in the experiment for five dimensions). Our approach based on hashing alone does not perform well as its execution time grows very rapidly. The problem with using hashing to compute distance calculations is that the precision degrades the bigger the  $\epsilon$  grows. Using the representative point approximation as well, however, again reduces the performance considerably, to the point where it again outperforms  $\rho$ -approximate DBSCAN.

For a very big  $\epsilon$  (which is rarely used in real clustering applications), however, the  $\rho$ -approximate DBSCAN outperforms ADvaNCE. This effect is due to the dataset rather than the algorithm. As outlined previously, the data points are normalized to an interval of  $[0, 10'000]$  in all dimensions. In these experiments, however, the  $\epsilon$  finally grows to 50'000 and, for example, in 5D the cell width is almost 23'000, resulting in only 5 cells in each dimension. With this few cells




Figure 5: Average execution time for experiments in 5, 7 and 9 dimensions with increasing dataset size.

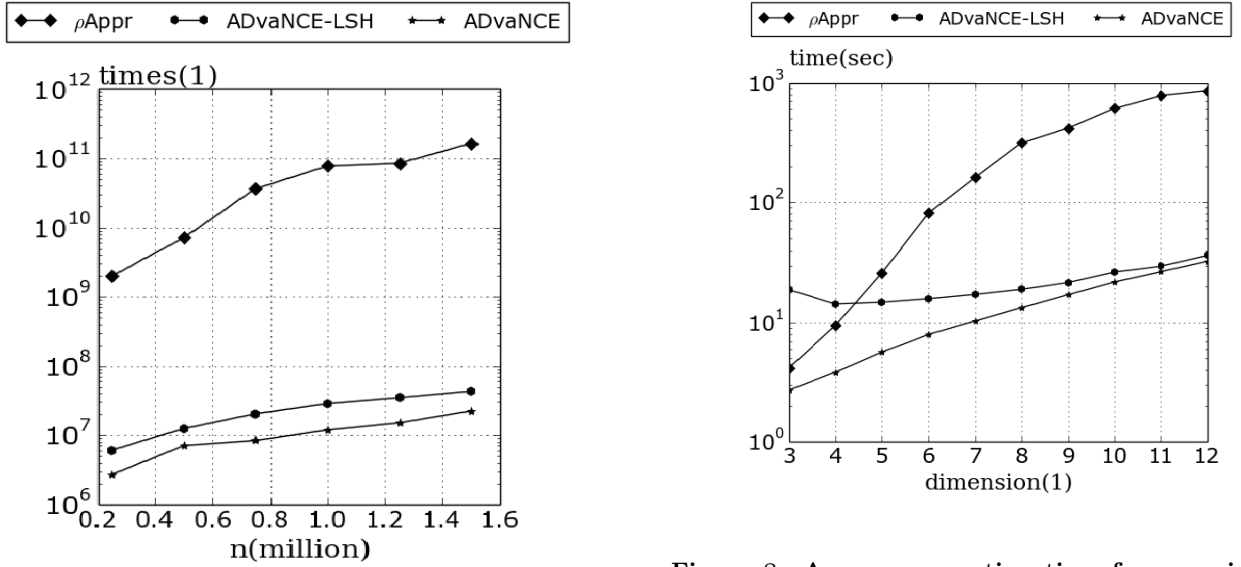


Figure 6: Number of distance calculations with increasing dataset size.

ADvANCE cannot considerably reduce and approximate the number of distance calculations.

#### 7.4.3 Increasing Dimensions

We also test the performance, i.e., execution time of our approach with an increasing number of dimensions (beyond 9 dimensions) and compare it with  $\rho$ -approximate DBSCAN. As we increase the dimension, we fix the remaining parameters: we use a synthetic dataset with three million data points and an  $\epsilon$  of 5'000.

Figure 8: Average execution time for experiments with an increasing number of dimensions.

As the result in Figure 8 shows, although concave down and appearing to converge, the execution time of  $\rho$ -approximate DBSCAN grows faster than the one of both ADvANCE versions, resulting in a speed up of more than one order of magnitude for the biggest dimension tested.

The higher the dimension,

Also here, the relative difference between the two versions of ADvANCE indicates that the higher the dimension, the more the hashing approximation contributes to the result as the execution time reduction thanks to the representative points approximation diminishes with increasing dimensionality. Towards the biggest number of dimensions tested, the execution time for  $\rho$ -approximate DBSCAN and our ap-






Figure 7: Average execution time for experiments in 5, 7 and 9 dimensions with increasing  $\epsilon$ .

proach are parallel with our approach still outperforming the state-of-the-art by more than one order of magnitude.

The results in Figure 9 show a similar picture. In this experiment we measure the number of distance calculations needed for an increasing number of dimensions for the same experimental setup. As the figure shows, the trend in both curves is the same as for the execution time in Figure 8 and consequently the number of distance calculations is indeed what drives the execution time.

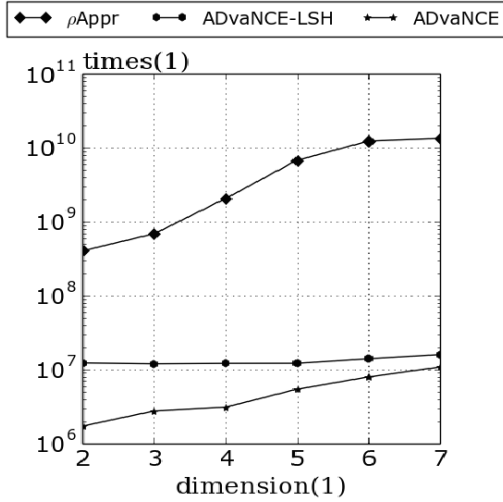


Figure 9: Number of distance calculations with an increasing number of dimensions.

## 7.5 Real World Datasets

As a final litmus test of the overall execution time of

ADvance, we also compare it against  $\rho$ -approximate DBSCAN on real world data. More specifically, we execute ADvance and  $\rho$ -approximate DBSCAN on the KDD cup '99 dataset as well as on the household dataset [4]. Given that the number of points per cell is very unbalanced in the real world datasets, we set  $maxNum$  for the representative approximation to  $\sqrt{A} \times M$  where  $M$  is the maximum number of points in a cell and  $A$  is the average number of points per cell.


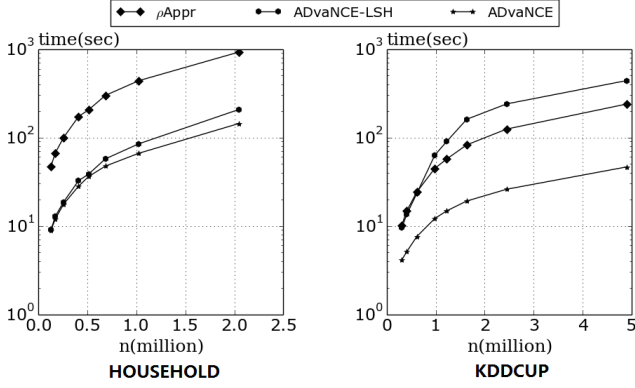


Figure 10: Average execution time for experiments on real data with increasing epsilon.

As the results in Figure 10 show, ADvance generally outperforms  $\rho$ -approximate DBSCAN. As we have already seen on the synthetic data, the ADvance version which also features the representative points approximation further improves the efficiency and in general improves execution time by a factor of  $5\times$ . Interestingly, for the KDD cup data the trend both versions of ADvance exhibit is

similar to  $\rho$ -approximate DBSCAN but ADvANCE is still one order of magnitude faster. For the household dataset both versions of ADvANCE scale substantially better compared to  $\rho$ -approximate DBSCAN: the execution time of  $\rho$ -approximate DBSCAN grows very fast while ADvANCE stays nearly flat.

In a next experiment we measure the execution time on the same real world datasets but we increase the size of the datasets. More precisely we start with a tenth of the dataset and then increase it until we arrive at their full sizes.



**Figure 11: Average execution time for experiments on real data with increasing dataset size.**

As the results in Figure 11 show, ADvANCE also outperforms  $\rho$ -approximate DBSCAN by one order of magnitude. The trends are similar, but for an increasing dataset size, ADvANCE has an edge over ADvANCE-LSH as the gap between the trends grows bigger.

## 8. IN-DEPTH ANALYSIS


Crucial to understanding the benefits of ADvANCE is to analyze in detail its behavior. In the following we thus first analyze the impact of increasing dimension and  $\epsilon$  on the number of points (and thus execution time) and further analyze ADvANCE through a breakdown of the execution time.

### 8.1 Number of Representative Points

As the dimension grows bigger, the maximum number of points and the average number of points in each cell is getting smaller (because size of the cell is getting smaller:  $cellWidth = \frac{\epsilon}{\sqrt{(d)}}$  and because the seed spreader has more directions to spread the data). Figure 12 confirms this by measuring the maximum as well as average number of points per cell for an increasing number of dimensions. Consequently, ADvANCE will perform well for an increasing number of dimensions even if only ADvANCE-LSH, the version which does not curb the number of points per cell, is used.

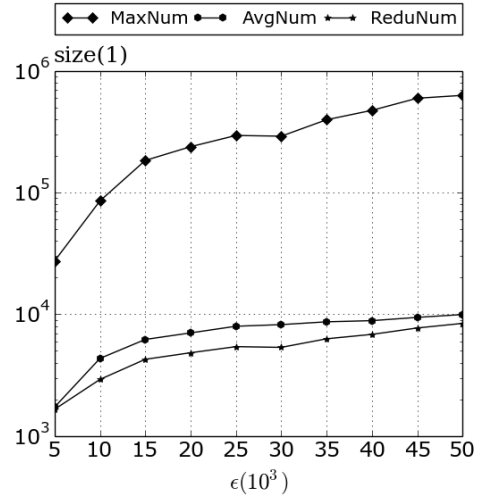
For an increasing  $\epsilon$  the picture is different as Figure 13 shows. In this experiment we measure the maximum (*Max Num*), the average (*Avg Num*) as well as the reduced number (*Redu Num*) of points ADvANCE considers. The reduced number of points is determined based on the maximum number of points in a cell  $M$  and  $minPts$  as follows:  $\sqrt{minPts \times M}$ .

The results clearly show that as  $\epsilon$  grows bigger, the size



**Figure 12: Maximum and average number of points per cell for increasing dimensions.**

of cell grows as well, and the number of points in each cell, however, is unbalanced: some cells contain a very large number of points while others contain only few. The second optimization of ADvANCE, representative points, effectively addresses this issue by reducing the number of points considered (while only sacrificing little accuracy).



**Figure 13: Maximum, average and reduced number of points per cell for an increasing  $\epsilon$ .**

### 8.2 Algorithm Breakdown

To understand the behavior of ADvANCE, we analyze the performance of its major building blocks *Determine Core Point LSH* and *Merge Clusters LSH* in the main iteration of our algorithm. We analyze the behavior for an increasing  $\epsilon$ . As the results in Figure 14 (left) shows, if we only use ADvANCE-LSH (i.e., ADvANCE without the representative

points approximation) as  $\epsilon$  grows, the time for determining core points will increase considerably from less than 0.1s to exceeding 100s. This is because in this step, for every non-core point, we have to search for all points in the approximate neighborhood to see if this point is a core point. As  $\epsilon$  grows larger, the number of points in the approximate neighborhood grows substantially and the time of this step will become the bottleneck. If we limit the number of points in each cell through the representative points approximation, however, the execution time is curbed.

As Figure 14 (right) shows, as  $\epsilon$  grows, if we only use hashing algorithm but not limit the number of points, the time for *Merge Clusters LSH* also increases. The increase, however, is moderate, going from less than 1 second to 10 seconds. This is because in this step we only find one core point in the approximate neighborhood and then merge cells for every point. As  $\epsilon$  grows bigger, the number of points in the approximate neighborhood also grows, but the time of this step consequently does not change substantially. If we limit the number of points in each cell to the representative number of points, the execution time of this step remains the same.

Assessing the execution time of these two functions, we see that if we only use ADvANCE-LSH, the execution time will increase as  $\epsilon$  grows, mainly because of the time spent on *Determine Core Points*. The execution time of both optimizations we use in ADvANCE, however, remains almost the same (growing slightly) for an increasing  $\epsilon$  thanks to the representative points approximation.

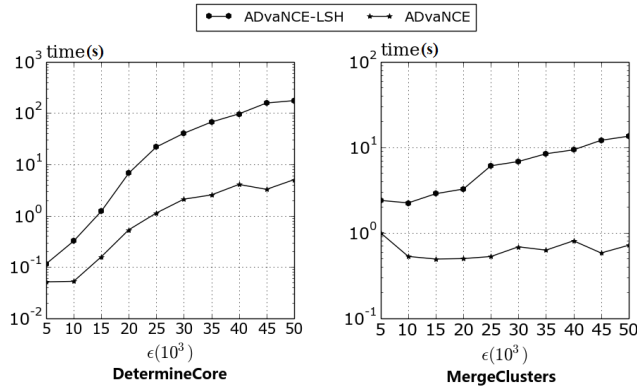


Figure 14: Breakdown of the execution time of ADvANCE.

## 9. RELATED WORK

In the discussion of related work we first give a brief overview of exact density-based clustering approaches before we move on to discuss in more detail the approximate algorithms devised.

### 9.1 Exact Approaches

Arguably the most renowned density based clustering approach is DBSCAN [11] itself. Given a set of points in typically high-dimensional space, it groups together points that are closely packed together (points with many nearby neighbors), marking as outliers points that lie alone in sparse regions. More formally, DBSCAN takes parameters  $\epsilon$  and

$minPts$  and identifies core points (points with at least  $minPts$  within radius  $\epsilon$ ). DBSCAN then assigns core points within distance  $\epsilon$  of each other (along with non-core points within distance  $\epsilon$ ) to a cluster. All other points are considered noise. DBSCAN can efficiently partition dense and sparse space but cannot deal well with varying density.

OPTICS [3, 7] addresses this issue of DBSCAN. Instead of using a fixed distance predicate,  $\epsilon$  is only used as a maximum distance and the distance is set in different areas of the dataset such that clusters can also be identified in areas of varying densities.

### 9.2 Approximate Approaches

More important to put our work in context, several approximate DBSCAN algorithms have also been devised in recent years.

ADBSCAN [18] approaches the problem of accelerating the clustering of DBSCAN by reducing the number of region queries needed. The approach is based on the assumption that clusters are discovered by executing range queries on the dataset. It consequently tries to minimize the number of range queries. To do so, it defines skeletal points as the minimum number of points where range queries (with radius  $\epsilon$ ) need to be executed to capture all clusters. The problem of finding the skeletal points is NP-complete and the ADBSCAN thus essentially proposes how to approximate the set of skeletal points using a genetic algorithm. By approximating the skeletal points, of course also the result of the clustering computation is approximate.

IDBSCAN [6] works similarly and approximates and thus accelerates DBSCAN through reducing the range or neighborhood queries needed. Instead of executing a query for all points in the vicinity of a core point, IDBSCAN only samples the neighborhood and executes a query on a subset of the points. To do so the neighborhood of a core point is partitioned into quadrants and a range query is only executed for the points in quadrants randomly chosen. The speedup achieved reaches a factor of six for the datasets tested. However, no guarantees about the results can be made.

l-DBSCAN [17] uses the concept of *leaders* to accelerate clustering. Leaders are a concise but approximate representation of the patterns in the dataset. l-DBSCAN on a first levels clusters the leaders instead of the full dataset thereby accelerating the process. In a second step it replaces the leaders by the actual points from the dataset. With this two level process l-DBSCAN outperforms the precise version of DBSCAN by a factor of at most two.

$\rho$ -approximate DBSCAN [12] guarantees the result of its clustering to be between the exact result of DBSCAN with  $(\epsilon, minPts)$  and  $(\epsilon \times (1 + \rho), minPts)$ . The approximation comes with a substantial speedup and scalability:  $\rho$ -approximate DBSCAN is proven to run in  $O(n)$ . The approach accomplishes this by using a grid approach where data points are assigned to a grid with cell width  $\frac{\epsilon}{\sqrt{d}}$ . For an exact result the algorithm has to connect/combine all pairs of cells  $c_1, c_2$  that (a) contain at least  $minPts$  and (b) contain two points ( $p_1 \in c_1$  and  $p_2 \in c_2$ ) within distance of each other  $\epsilon$ . This problem is known as the biochromatic closest pair (BCP) and solving it precisely is exceedingly costly. However, solving BCP approximate (in  $\epsilon \times (1 + \rho)$ ) can be accomplished in linear time. Clearly, the main benefit of this approach are the proven and firm theoretical bounds on error/approximation as well as the execution time.

Pardicle [15] approximates DBSCAN in a supercomputing environment. It accelerates the basic DBSCAN approach through parallelization: after partitioning the multi-dimensional dataset into contiguous chunks, DBSCAN finds clusters in each chunk in parallel on different CPUs/cores of the supercomputing infrastructure. To account for clusters that span several chunks (analyzed independently and in parallel on different cores), Pardicle replicates the border (of width  $\epsilon$ ) of a chunk and copies it to adjacent chunks, i.e., to cores computing clusters. Doing so in essence makes the chunks overlap but also ensures correctness of the result. To reduce the overhead in terms of replication and distance calculations not the exact border is calculated but is sampled. Doing so approximates the results: the result in each chunk is accurate, but merging of clusters between adjacent chunks may fail.

## 10. CONCLUSIONS

In this paper, we develop ADvaNCE, a novel approach for approximating density-based clustering. ADvaNCE builds on a Grid-based optimization of DBSCAN and uses two approximations to accelerate and enable scalability of the clustering process. Understanding that distance calculations are the major contributor to the execution time of DBSCAN, ADvaNCE accelerates clustering by (a) approximating distance calculations and (b) reducing the distance calculations required, thereby further approximating the result. To achieve the necessary level of approximation, ADvaNCE iterates over intermediate results to monotonically increase the accuracy.

Approximate clustering results oftentimes are sufficient as has been established by related work [3, 12]. Compared to the most recent state-of-the-art algorithms in approximate density-based clustering [12], ADvaNCE achieves the same accuracy but manages to do so more than one order of magnitude faster (30x) in the best case. With our experimental evaluation we also show that the ADvaNCE executes faster when varying parameters like  $\epsilon$ , dataset size or the number of dimensions. In general, ADvaNCE scales better than the state-of-the-art algorithms.

## 11. REFERENCES

- [1] S. Adaszewski, J. Dukart, F. Kherif, R. Frackowiak, and B. Draganski. How early can we predict Alzheimer's disease using computational anatomy? *Neurobiology of Aging*, 34(12):2815 – 2826, 2013.
- [2] A. Andoni and P. Indyk. Near-optimal Hashing Algorithms for Approximate Nearest Neighbor in High Dimensions. *Communications of the ACM*, 51(1), 2008.
- [3] M. Ankerst, M. M. Breunig, H.-P. Kriegel, and J. Sander. Optics: Ordering points to identify the clustering structure. In *Proceedings of the 1999 ACM SIGMOD International Conference on Management of Data*, SIGMOD '99.
- [4] K. Bache and M. Lichman. UCI Machine Learning Repository, 2013.
- [5] J. L. Bentley. Multidimensional Binary Search Trees Used for Associative Searching. *Communications of the ACM*, 18(9), 1975.
- [6] B. Borah and D. Bhattacharyya. An Improved Sampling-based DBSCAN for Large Spatial Databases. In *Proceedings of International Conference on Intelligent Sensing and Information Processing*, 2004.
- [7] M. M. Breunig, H.-P. Kriegel, R. T. Ng, and J. Sander. Optics-of: Identifying local outliers. In *Principles of Data Mining and Knowledge Discovery*, volume 1704 of *Lecture Notes in Computer Science*. 1999.
- [8] M.-S. Chen, J. Han, and P. Yu. Data mining: an Overview from a Database Perspective. *IEEE Transactions on Knowledge and Data Engineering*, 8(6), 1996.
- [9] L. M. Collins and C. W. Dent. Omega: A general formulation of the rand index of cluster recovery suitable for non-disjoint solutions. *Multivariate Behavioral Research*, 23(2), 1988.
- [10] M. Datar, N. Immorlica, P. Indyk, and V. S. Mirrokni. Locality-sensitive Hashing Scheme Based on P-stable Distributions. In *Proceedings of the Twentieth Annual Symposium on Computational Geometry*, SCG '04, 2004.
- [11] M. Ester, H.-P. Kriegel, J. Sander, and X. Xu. A density-based algorithm for discovering clusters in large spatial databases with noise. In *Proceedings of the 2nd International Conference on Knowledge Discovery and Data Mining*, 1996.
- [12] J. Gan and Y. Tao. DBSCAN Revisited: Mis-Claim, Un-Fixability, and Approximation. In *Proceedings of the 2015 ACM SIGMOD International Conference on Management of Data*, SIGMOD '15.
- [13] A. Gunawan. A faster algorithm for DBSCAN. Master's thesis, Technical University of Eindhoven, March 2013.
- [14] G. Murray, G. Carenini, and R. Ng. Using the Omega Index for Evaluating Abstractive Community Detection. In *Proceedings of Workshop on Evaluation Metrics and System Comparison for Automatic Summarization*, 2012.
- [15] M. M. A. Patwary, N. Satish, N. Sundaram, F. Manne, S. Habib, and P. Dubey. Pardicle: Parallel Approximate Density-based Clustering. In *Proceedings of the International Conference for High Performance Computing, Networking, Storage and Analysis*, SC '14.
- [16] T. Venetis, A. Ailamaki, T. Heinis, M. Karpathiotakis, F. Kherif, A. Mitelpunkt, and V. Vassalos. Towards the Identification of Disease Signatures. In *Proceedings of the 8th International Conference on Brain Informatics and Health BIH*, 2015.
- [17] P. Viswanath and R. Pinkesh. l-DBSCAN : A Fast Hybrid Density Based Clustering Method. In *Proceedings of the International Conference on Pattern Recognition*, 2006.
- [18] S. Yeganeh, J. Habibi, H. Abolhassani, M. Tehrani, and J. Esmaelnezhad. An Approximation Algorithm for Finding Skeletal Points for Density Based Clustering Approaches. In *Proceedings of the IEEE Symposium on Computational Intelligence and Data Mining*, 2009.